

# Master-Thesis Colloquium

A Hardware-Supported Boot and Debug Unit for RISC-V + eFPGA SoCs

Stephan Proß

March 16, 2023

# Motivation

## FABulous Framework

- Embedded FPGA generation
- Open Source
- Modular & customizable
- Target: ASIC/FPGA
- Bit-Stream generation
- Case Study: RISC-V + eFPGA

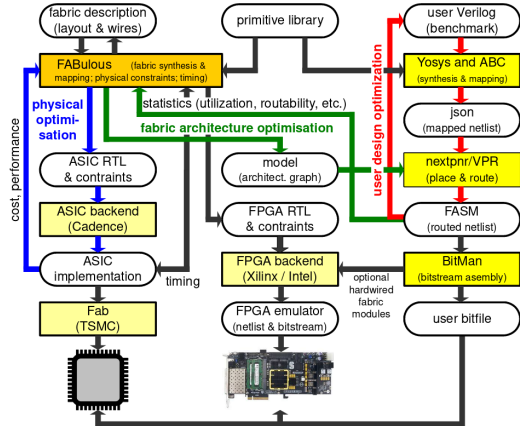


Figure: FABulous Framework[3]

# Combining RISC-V & eFPGA

## Case Study

RISC-V + eFPGA hybrid:

- 32-bit RISC-V CPU
- eFPGA as custom instruction extension

⇒ No independence

## Independence Benefits

- Application versatility
  - Cooperative, Boss-Worker, Parallel
- Robustness




Figure: RISC-V + eFPGA hybrid[3]

# FABulous SoC

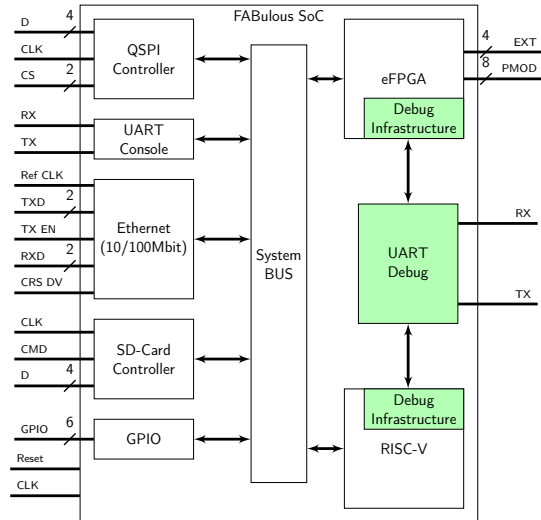
## FABulous SoC

CVA6 RISC-V + eFPGA:

- 32-bit RISC-V CPU
- Linux support
- Independent eFPGA

Goals:

- Academic usage
- Versatility
- Traceability



# Design Goals

## Traceability

- Single step debugging of CPU
- Debugging from reset
- System bus activities
- Memory access

## Versatility / Utility

- Features exceeding debug

## Robustness

- Independence from system state
- Re-use of existing solutions

How to bring up CPU & FPGA ?

⇒ Anatomy of a RISC-V CPU

# RISC-V CPU

## RISC-V ISA

- Extensions
  - 32-bit Base Integer
  - FPU, Vector ...

## CVA6 Implementation

- Design Goal: OS-Support (Linux)
- Single Core, Single Hart
- Customizable
  - Data width (32,64)
  - FPU, MMU ...

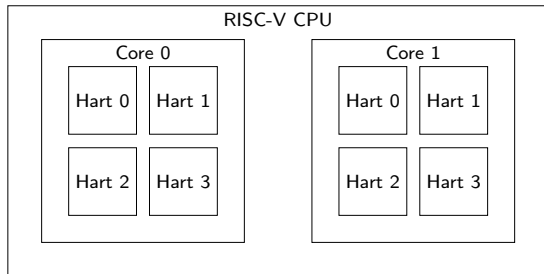


Figure: RISC-V CPU Example

# Booting RISC-V

## Power-On/Reset

### Hard-coded Preliminary Boot-Code

- Harts begin at Reset-Vector (0x100)
- Setup of registers:
  - Hart-ID
  - Flattened-Device-Tree (FDT) address
  - Start-Address

## Boot-Code

### Resides in (Flash-)Memory

- Setup of hardware using FDT, application environment
- Handover to
  - Bare-metal program
  - OS

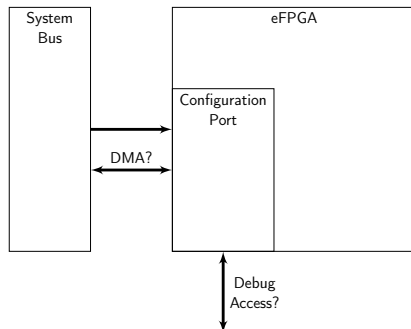
## Requirements

Access to system (flash-)memory

# Booting eFPGA

## Power-On/Reset

- Loading configuration
  - External write (CPU, debug infrastructure)
  - DMA from memory (flash, RAM)
- Fabric enable



## Requirements

Configuration port access



# Debugging Embedded Systems

## General Debugging

- R/W access
  - Process memory, registers
- Control of execution state
  - Halt, resume
- Code execution / patching
  - Breakpoints

Requirement: **System Calls**

- Unix: ptrace

## Issue: Embedded Systems

Without OS:

- System calls unavailable
- Remote target

## Requirements

- Remote communication
- Hardware debug support

# Debugging RISC-V

## Debugger GDB

- Remote debug using TCP/IP

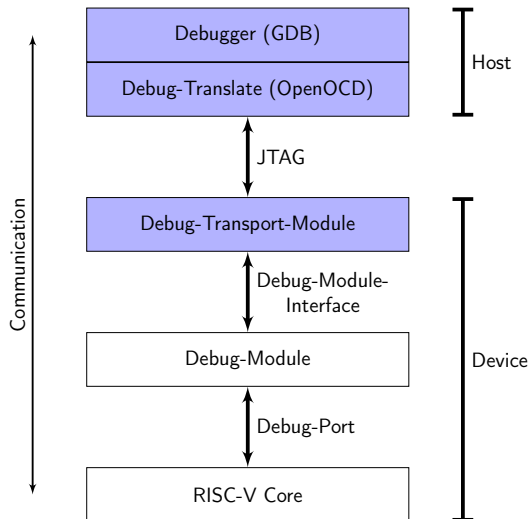
## Debug Translation

### Open On-Chip Debugger (OpenOCD)

- GDB interface (gdbserver)
- Translates to target (RISCV, ARM, ...)
- Handles Transport (JTAG, SPI..)

## Debug Transport Module

### Access to Debug-Module-Interface



# Debug Module

## Debug Module

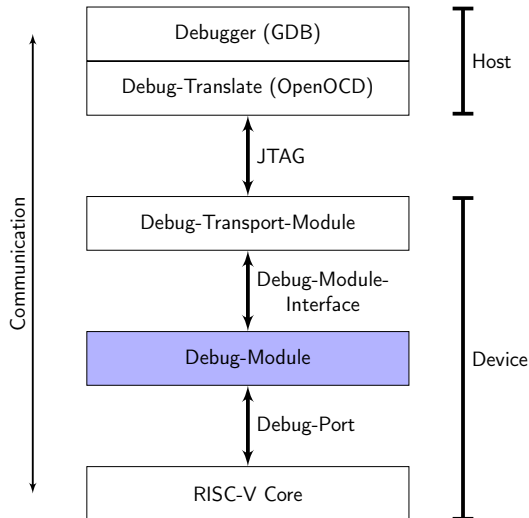
Implements:

- Reset, halt, resume of Harts
- Arbitrary command execution
- Register & memory access

## Implementation Options

One of:

- Abstract commands
- Program buffer
- System bus access



# Existing Infrastructure

## RISC-V Debug Project

Integrated into CVA6

Features:

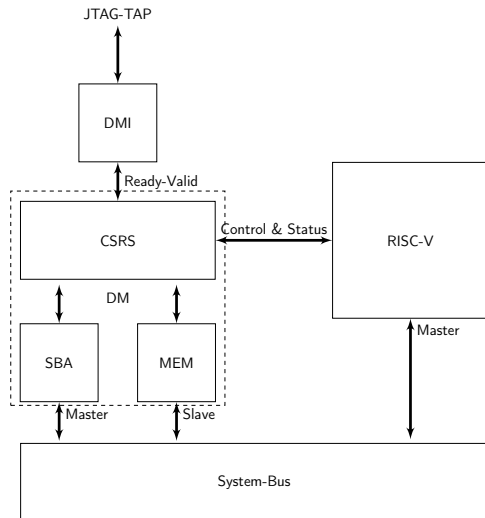
- Register access
- System bus access
- JTAG as transport

## Debug Module

**CSRS** Debug registers

**MEM** Program buffer

**SBA** System bus access



# Remaining Work

## OpenOCD

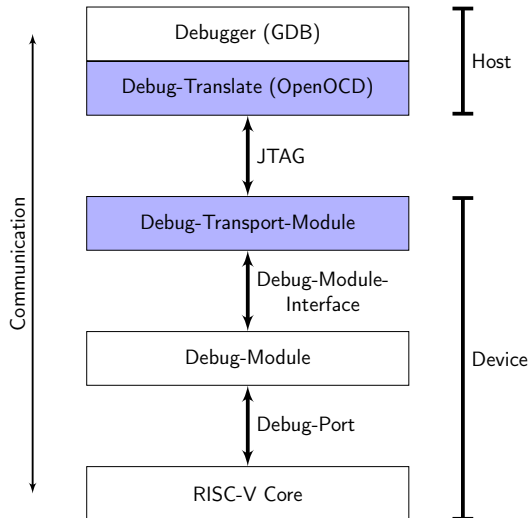
Available transports

- JTAG, SPI ...

No UART support!

## Debug-Transport-Module

- JTAG only
- Additional features



# JTAG vs. UART

## JTAG

### Wires:

- **TDI**, **TDO**: Serial Data I/O
- **TMS**: State Machine Control
- **TCK**: Clock from host

### Properties:

- Synchronous
- Instruction (IR) and data (DR) registers
- **TMS** can trigger reset

## UART

### Wires:

- **RX**, **TX**: Serial Data I/O

### Properties:

- Asynchronous
- Transmission in frames:
  - 1 Start (1 bit)
  - 2 Data (5-9 bits)
  - 3 Parity (0-2 bits)
  - 4 Stop (1-2 bits)

# UART-Integration

## JTAG-UART Adapter

Host software:

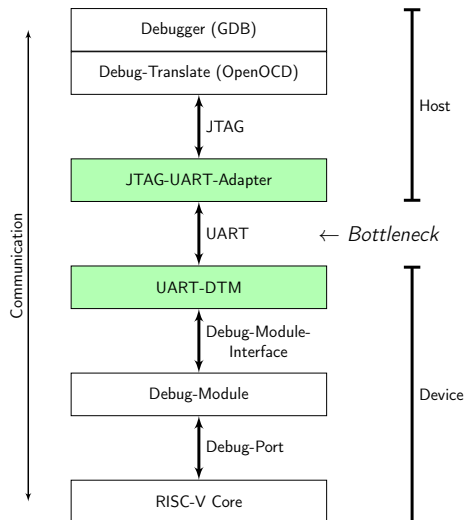
- Maintainability
- Speed

Purpose:


- Functional translation

## UART DTM

- UART-Interface
- (Legacy) features



# JTAG-UART Adapter



## Components


- JTAG-FSM
  - (De-)Serialization
- UART-Interface
  - Byte-wise exchange
- Adapter
  - Functional translation

## Issue

Distinction between data and instruction



# UART Communication



Common Case 8-bits Data, 0-bits Parity, 1-bit Stop

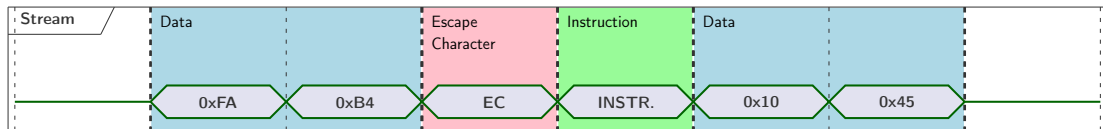
## Data Rate

- Symbol/Baud rate  $f_s$
- Serial  $\Rightarrow f_s = \text{raw bitrate } f_b$
- Common Case: effective data rate  $f_{\text{eff}} = f_s \text{ 80\%}$

## Debug Protocol

- Mark instruction/data
- Low overhead

# Protocol: Instruction Escape



## Protocol

- Escape Character (EC) before instr.
- EC as data  $\Rightarrow$  repeat

## Properties

- $40\% f_s \leq f_{eff} \leq 80\% f_s$
- Lin. Dist. :  $p = 1/256$ 
  - $f_{eff} = 79.7\% f_s$

Break even point:  $p = 1/7 \approx 37/256$

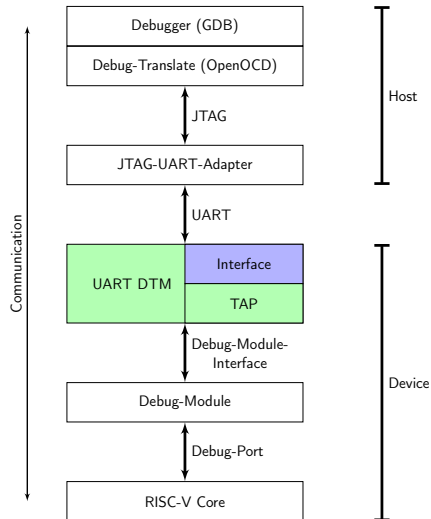
# UART-DTM

## UART-Interface

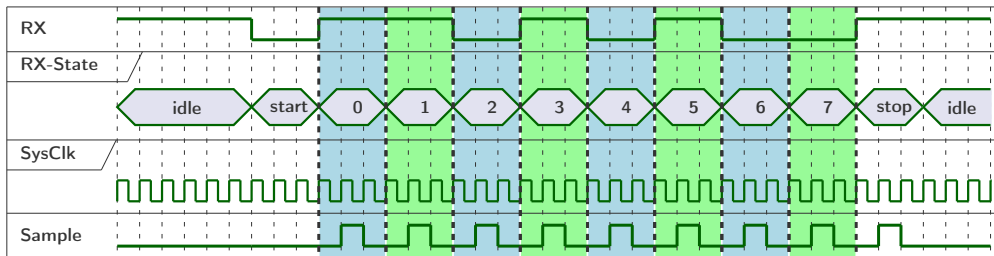
- Sampling of RX
- Un-/packing data from/to frames

## UART-TAP

- Debug-Protocol interpretation
- Read/write registers



# UART Interface



## RX Sampling

- Process start w. falling edge
- Centered in symbol period


## Sampling Interval

- $f_{sys} = 10MHz$
- $f_s = 1MBd$

Clock divider:

- $C := \lfloor f_{sys}/f_s \rfloor = 3$

# Sampling Interval



## Sampling Interval

- $f_{sys} = 10MHz$
- $f_s = 3MBd$

Clock divider:

- $C := \lfloor f_{sys}/f_s \rfloor = 3$

## Issue

Integer Division of  $f_{sys}$   $f_s$   
⇒ Remainder causes drift!

# Sampling Solutions

## 1. Solution: Correction w. Edge Detection

Edges in data used as synchronization

- Data dependency!
- Expensive logic

⇒ Only Mitigation

## 2. Solution: Remainder Counting

“Add” remainder to delay Sampling.

- Data Independent
- Small additional counter

⇒ Actual Prevention

## Issue

Target device might be too slow / busy

## Solution

Software Signaling:

- FIFO-Buffer w. fill-level signal
  - Latency is critical
- Halt-Signal to host
  - ASCII: XON/XOFF?
  - Escape Character!

## Optimal Buffer Depth ?

- EC transmission time
- TX can be busy
- Host Latency?

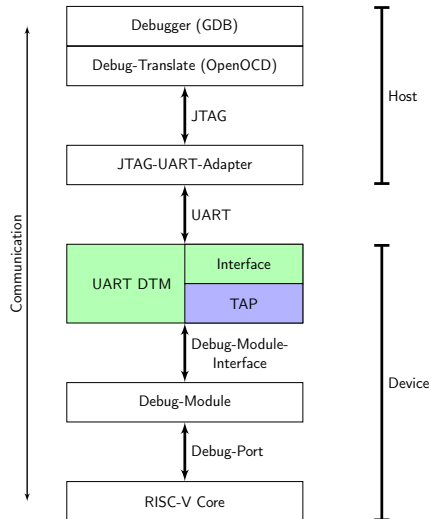
# UART-TAP

## Requirements

- Legacy Registers
- Peripheral access
- Protocol Interpretation
- R/W Streaming

## Goals

- Resource Efficiency
- Utility
- Extendability





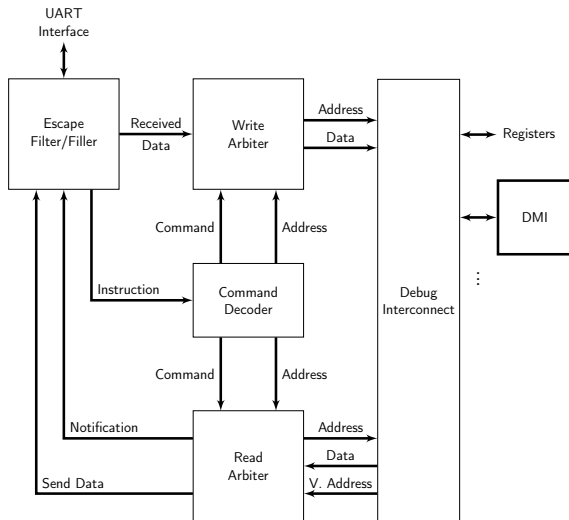
# TAP Architecture

## Instructions

- Continuous Write
- Single / Continuous Read
- Reset

## Additional Features

- Escape Character handling
- Byte-wise (de-)serialization
- Bi-directional streaming
- Host-Notification



# Debugging FPGAs

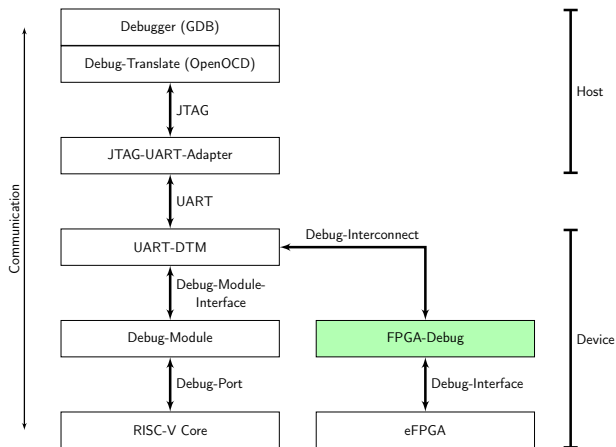
## Debug Stack

- JTAG replaced with UART
- Original functionality preserved
- Efficient streaming

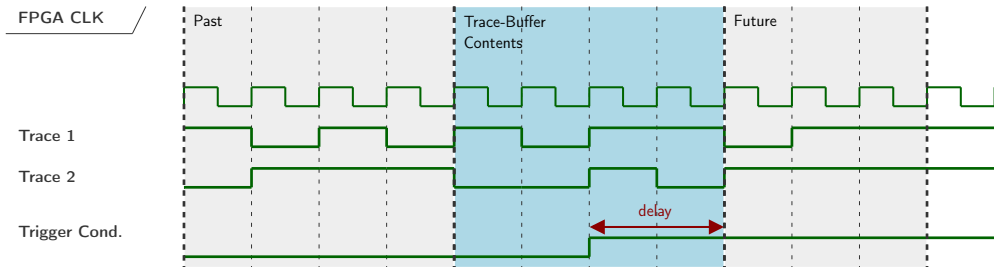
## FPGA Debugging

### Common Capabilities:

- Trace Capture
- Trigger Logic
- Basic data exchange



# Trace Capture



## Trace Buffer

- Ring Buffer
  - Continuous Signal Capture
  - Trigger cond. stops capture
- ⇒ Logic Analyzer

## Trigger Condition

- Arbitrarily complex
- Delay:
  - Ratio of history to pre-history

# FPGA Debug Infrastructure

## Why?

Debug infrastructure in soft logic:

- Less logic for user design

Large area overhead:

- 1x 1kB SRAM Cell  $\approx 32$ x LUT4

⇒ hard logic smaller

## Proposed Solution

Streaming Trace Buffer (STB)




Figure: Caravel with eFPGA[2][1]

# Streaming Trace Buffer

## Components

- 1kB SRAM cell used as:
  - Trace-Buffer
  - RW-Buffer
- FPGA-Interface
  - Trace-Capture
  - Clock Domain Crossing
- System-Interface
  - Connects to debug interconnect

## Modes

- Trace Mode
  - Variable number of trace signals
  - More signals less history
- Stream Mode
  - Bi-directional

## Trigger Logic?

- Relegated to user logic
- Single signal is trigger
- Variable delay

# Overview

## FPGA Interface


- Trace-Capture
  - De-/Serialization
- Trigger Delay

## Memory Controller

- Pointer Handling
- RW Strobing

## System Interface

- Control & Status
- 32-bit data exchange



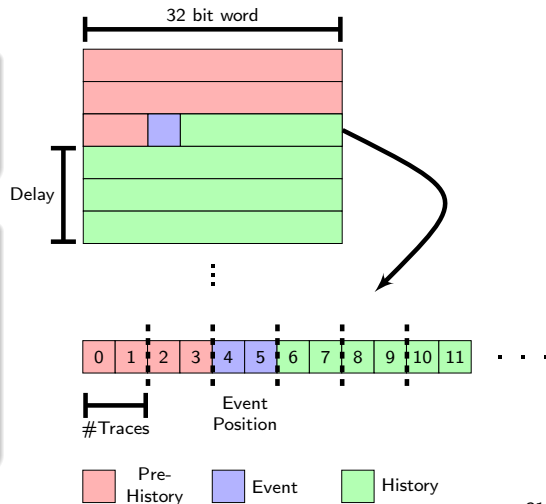
# History Reconstruction

## Trigger event:

- Full SRAM readout
  - First value = history begin

## Required Information

- Control Registers:
  - Delay  $\Rightarrow$  Memory word in readout
  - Number of Traces  $\Rightarrow$  Granularity
- Status:
  - Event Position Status



# Implementation

## FPGA Emulation

Xilinx Vivado 2021.1 targeting Nexys-4 DDR Board

- System Clock : 25 MHz
- Baud Rate : 3 MBd

## ASIC

- SRAM-Macros: OpenHW Skywater 130nm
- Tool-Chain: Yosys
- Available space: 10mm<sup>2</sup>

Area Requirements?



# Area STB

Component	Area ( $\mu\text{m}^2$ )	Total%	LUTs	Registers
STB	184204.5	100.0	277	340
- FPGA-Interface	12406.6	6.7	248	306
- Memory-Controller	171632.5	93.2	45	76
- - SRAM-Cell	167999.0	91.2		
- System-Interface	217.8	0.1	4	6
STB w.o. SRAM-Cell	16205.5	8.8	277	340

## Observation

- SRAM-Cell dominates cost
  - Surrounding logic fractional

Total Area?

# Debug Logic Area

Component	LUTs	LUT%	Registers	Register%
STB w.o. SRAM-Cell	277	100.0	340	100.0
UART-DTM	331	119.5	401	117.9
Debug Module	455	164.3	676	198.8
Total	1063	383.8	1417	416.8
CVA6-CPU (32bit)	13529		7402	

## Observation

FPGA emulation:

- Costs of DTM & DM comparable to STB

## Approximation

- LUTs & registers  $\sim$  area
- $\Rightarrow$  Total Area  $\approx$  40% of SRAM-Cell

# Task Review

## Device-to-Host using UART

- UART-DTM ☒
- Host Software ☒
- Debug Protocol ☒

## FPGA Infrastructure


- Trace Capture ☒
- Data Exchange ☒
- Configuration Access (☒)  
⇒ Access via System-Bus

# Questions

# References I

- [1] *Chip Gallery — FABulous Dokumentation 0.1 Documentation*. URL: <https://fabulous.readthedocs.io/en/latest/gallery/index.html> (visited on 03/15/2023).
- [2] *GitHub - Nguyendao-Uom/eFPGA\_v3\_caravel*: <https://caravel-user-project.readthedocs.io>. URL: [https://github.com/nguyendao-uom/eFPGA\\_v3\\_caravel](https://github.com/nguyendao-uom/eFPGA_v3_caravel) (visited on 03/15/2023).
- [3] Dirk Koch et al. “FABulous: An Embedded FPGA Framework”. In: *The 2021 ACM/SIGDA International Symposium on Field-Programmable Gate Arrays*. FPGA '21. New York, NY, USA: Association for Computing Machinery, Feb. 2021, pp. 45–56. ISBN: 978-1-4503-8218-2. DOI: 10.1145/3431920.3439302. URL: <https://doi.org/10.1145/3431920.3439302> (visited on 01/13/2023).

# Protocol: JTAG Over UART




## Protocol

- Wrap **TDI**, **TMS** in a frame
- TCK generated by receiver

## Properties

- $f_{eff} = f_s$  40%
- JTAG-TAP reusable

# Protocol: Instruction Bit



## Protocol

- Reserve one bit of data
- Instruction: Address & command


## Properties

- $f_{eff} = f_s$  70%

Worse for 32-bit data:

- 5 Frames  $\Rightarrow f_{eff} = f_s$  64%

# Protocol: Instruction Packets



## Protocol

**Packet** Header, instruction,  
(length), data

## Properties

Assuming 4 bytes data:

- $f_{eff} = 53\% f_s$
- Break even point: 14 bytes data



# Additional Features

## Observations

- USB-to-UART chip single channel
- Additional data paths desirable

Direct RX,TX access to eFPGA:

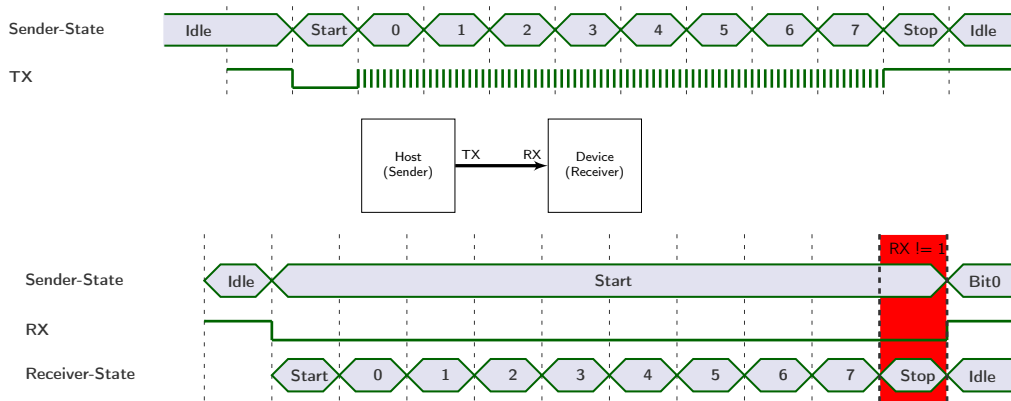
- Simple messaging
- UART-Interface in eFPGA

## Multi-Channel Support

Allow debug infrastructure and eFPGA usage of interface

- Auto detection
- Low protocol overhead

# Multi-Channel



## Condition

$$f_{Receiver} \geq 10 * f_{Sender}$$


# System Side

## Modular RV-Register

- Control & Status Interface
- Data Interface
  - 32 bit asynchronous R/W

## Control & Status Interface

- Trigger Delay
- Number of Traces
- STB Mode
- Trigger Event
  - Presence
  - Bit Position

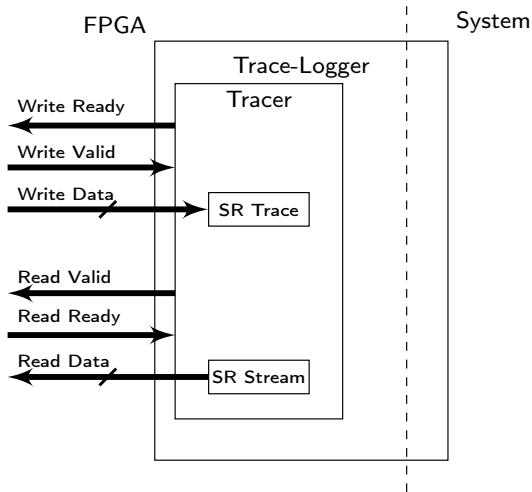


# Stream Mode

## Properties

- Bi-directional, async. R/W
- Ready-Valid signaling:
  - Trigger → Write Valid
  - Delayed Trg. → Read Ready
- Requirement: Matching I/O rates


⇒ Uni-directional modes



# Trace Mode

## Properties

- Traces continuously captured in memory
- Oldest traces output via Stream
  - $\Rightarrow$  STB daisy-chaining
- Trigger halts capture after delay



# STB Chaining

## Dual STBs


Allow:

- Tracing more signals in parallel
- Increase history depth
  - Chaining

## Chaining

Stream outputs oldest trace:

- Second STB captures trace



## Analysis

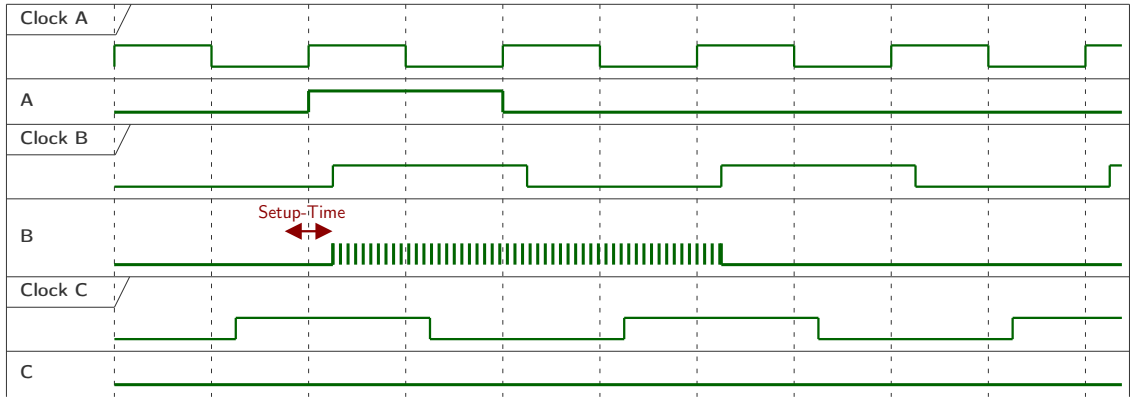
Protocol:

- $f_{eff} \approx 3\text{MBd} * 80\% = 2.4 \text{ Mb/s}$

Hardware Bottlenecks?

- UART-DTM ×
- RV-Register ×
- FPGA-Interface

# Clock Domain Crossing Issue



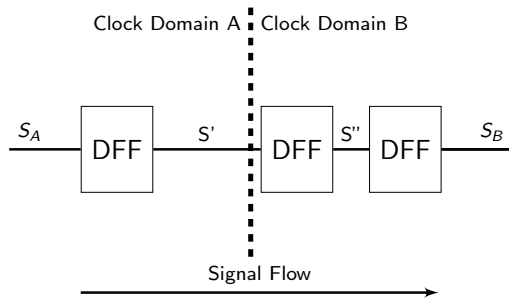


## Open-Loop Synchronizer

- Basis: FF in source and target domain
- Adds FF to path in target domain
- Eventually Accurate

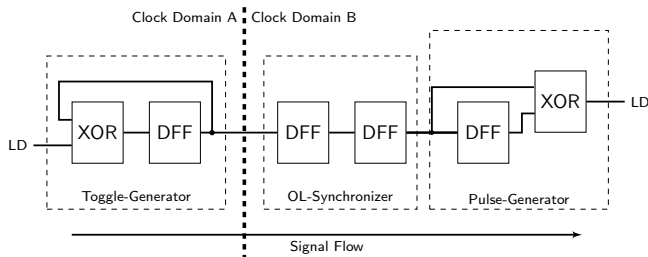
### Stability requirement:

- Source signal stable for 1.5 times the target period.



## Multi-Cycle-Path Formulation

- Multiple data signals, one synchronizer
- Data registered in both domains
- Sample Accurate
- Sync: Pulse  $\Rightarrow$  Toggle  $\Rightarrow$  Pulse

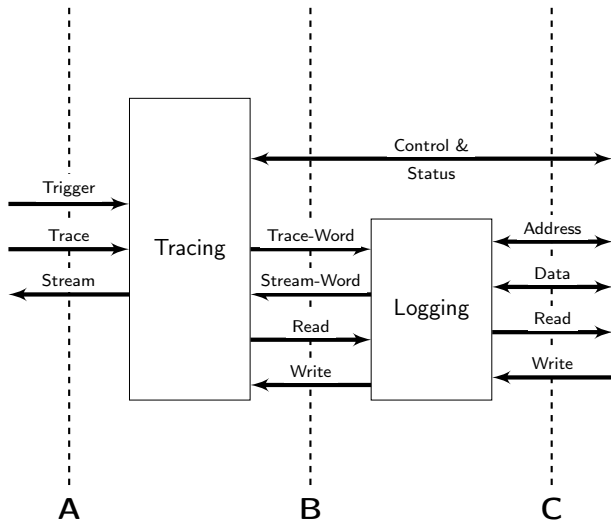


# CDC Design

## CDC Techniques

- Resource Costs
- Delay

How to partition logic into  
Clock Domains?



# Implementation FPGA-Interface

